

**Faculdade de Engenharia da Universidade do Porto**



# **Performance Evaluation of Single Core and Multi-Core Implementations**

**First Practical Project Report**

**CPD 2024/2024 - LEIC**

**Authors:**

Domingos Neto - up202108728@up.pt

Eduardo Baltazar - up202206313@up.pt

Joana Noites - up202206284@up.pt

|   |           |
|---|-----------|
| <b>Introduction</b>   | <b>3</b>  |
| <b>Problem Description</b>  | <b>3</b>  |
| <b>Algorithms</b>   | <b>3</b>  |
| Simple Multiplication   | 3         |
| Line Multiplication   | 4         |
| Block Multiplication  | 4         |
| Line Multiplication with parallel threads - First Solution  | 4         |
| Line Multiplication with parallel threads - Second Solution                                       | 4         |
| <b>Performance Metrics</b>  | <b>4</b>  |
| <b>Results and Analysis</b>   | <b>5</b>  |
| Simple Multiplication   | 5         |
| Line Multiplication   | 5         |
| Block Multiplication  | 6         |
| Line Multiplication with parallel threads - First Solution  | 6         |
| Line Multiplication with parallel threads - Second Solution                                       | 7         |
| <b>Conclusions</b>  | <b>7</b>  |
| <b>Annexes</b>  | <b>8</b>  |
| <b>A1 - Code Snippets</b>   | <b>8</b>  |
| A1.1 - Simple Multiplication  | 8         |
| A1.2 - Line Multiplication  | 8         |
| A1.3 - Block Multiplication   | 8         |
| A1.4 - Line Multiplication with parallel threads - First Solution                                 | 9         |
| A1.5 - Line Multiplication with parallel threads - Second Solution                                | 9         |
| <b>A2 - Simple Multiplication - Measurements and Graphs</b>                                       | <b>10</b> |
| A2.1 - Measurements Table   | 10        |
| A2.2 - C++ vs Java Execution Time vs Matrix Size  | 10        |
| <b>A3 - Line Multiplication - Measurements and Graphs</b>   | <b>11</b> |
| A3.1 - Measurements Table   | 11        |
| A3.2 - C++ Simple vs. Line Multiplication: Time Comparison  | 11        |
| A3.3 - C++ Simple vs. Line Multiplication: Cache Misses Comparison                                | 12        |
| A3.4 - C++ vs. Java Line Multiplication: Time Comparison  | 12        |
| <b>A4 - Block Multiplication - Measurements and Graphs</b>  | <b>13</b> |
| A4.1 - Measurements Table   | 13        |
| A4.2 - C++ Line vs. Block Multiplication - Time Comparison  | 13        |
| A4.3 - C++ Line vs. Block Multiplication - Cache Misses Comparison                                | 14        |
| A4.4 - C++ Simple vs. Line vs. Block Multiplication - Time Comparison                             | 14        |
| <b>A5 - Line Multiplication with Parallel Threads - First Solution - Measurements and Graphs</b>  | <b>15</b> |
| A5.1 - Measurements Table   | 15        |
| A5.2 - Line vs Parallel Line 1: Time Comparison   | 15        |
| A5.3 - Line vs Parallel Line 1: Cache Misses Comparison   | 16        |
| <b>A6 - Line Multiplication with Parallel Threads - Second Solution - Measurements and Graphs</b> | <b>16</b> |
| A6.1 - Measurements Table   | 16        |
| A6.2 - Line vs. Parallel Line 2: Time Comparison  | 17        |
| A6.3 - Line vs. Parallel Line 2: Cache Misses Comparison  | 17        |
| A6.4 - Parallel 1 vs. Parallel 2 - Time Comparison  | 18        |
| A6.5 - Parallel 1 vs. Parallel 2 - Cache Misses Comparison  | 19        |
| A6.6 - Parallel 1 vs. Parallel 2 - MFLOPS Comparison  | 19        |
| A6.7 - Parallel 1 vs. Parallel 2 - Efficiency Comparison  | 20        |

## Introduction

This work was developed as part of the curricular unit of CPD - Computação Paralela e Distribuída - and was based on the teachings of this unit.

In this report, we will analyze the effect on the processor performance of the memory hierarchy when accessing large amounts of data.

## Problem Description

In this project, the main objective is to analyze the performance of a processor during the execution of a matrix multiplication operation. Performance metrics, such as execution time and data cache misses for Level 1 (L1) and Level 2 (L2) caches, will be evaluated to assess single-core performance. These metrics will also be extended to a multi-core evaluation, with the addition of measuring the millions of floating-point operations per second (MFLOPS) to further measure computational efficiency. By examining these factors, the aim is to study how cache utilization and parallel processing impact overall system performance.

## Algorithms

To evaluate single-threaded performance, three distinct matrix multiplication algorithms were implemented and tested. The first algorithm (**Simple Multiplication**), provided in the work statement, performs a standard matrix multiplication. The second algorithm (**Line Multiplication**) multiplies an element from the first matrix with the corresponding row of the second matrix. The third algorithm (**Block Multiplication**) introduces a block-oriented approach, where the matrices are divided into smaller sub-matrices, and the multiplication is performed on them. **Simple Multiplication** and **Line Multiplication** were also implemented in Java so as to compare the performance of different languages for the same algorithm.

For multi-threaded performance evaluation, the **Line Multiplication** algorithm was modified according to two different solutions provided to us. The first solution parallelizes the outermost *for* loop, while the second one does the same to the innermost *for* loop. These modifications allow us to analyze how the placement of parallelism impacts the algorithm's performance

### *Simple Multiplication*

As mentioned earlier, this algorithm, provided in the work statement, iterates over the rows of the first matrix and multiplies them by the corresponding columns of the second matrix to compute the resulting matrix product. (Refer to Annex A1.1)

### *Line Multiplication*

This algorithm iterates over the elements of the first matrix and multiplies each element with the corresponding row of the second matrix, accumulating the results to produce the final matrix product. (Annex A1.2)

### *Block Multiplication*

This algorithm divides the matrices into smaller blocks of size *bkSize* x *bkSize*. It iterates over these blocks, multiplying corresponding blocks from the two matrices and storing the results into the output matrix.(Annex A1.3)

### *Line Multiplication with parallel threads - First Solution*

In this solution OpenMP's `#pragma omp parallel for` is used on the outermost for loop, this way, each thread independently processes a portion of the iterations of the *i* loop, while the inner loops execute sequentially within each thread. (Annex A1.4)

### *Line Multiplication with parallel threads - Second Solution*

In this solution OpenMP's `#pragma omp parallel for` is used on the innermost *for* loop, this way, each thread processes a portion of the *j* loop iterations independently, while the outer loops (*i* and *k*) run sequentially. (Annex A1.5)

## **Performance Metrics**

To evaluate the performance of the implemented matrix multiplication algorithms, several metrics were used. The main and most illustrative metric is **execution time**, which was measured for every algorithm considering differently sized matrices. This provided a direct comparison of the computational efficiency of every algorithm.

In addition to the execution time, performance counters such as the **Level 1 and Level 2 Data Cache Misses** (referred as L1 DCM and L2 DCM, respectively) were also accounted for, they were measured using the Performance API tool (PAPI). Measuring Data Cache Misses allows us to evaluate how efficiently each algorithm accesses the memory.

For the block-oriented algorithm, the impact of **block size** on performance was analyzed by testing different block sizes.

For the parallel implementations, **speedup** and **efficiency** were calculated to quantify the performance gains achieved through parallelization. Speedup was computed as the ratio of the sequential execution time to the parallel execution time. Efficiency, on the other hand, was derived by dividing the speedup by the number of cores used, which in our case was 8. Additionally, the computational throughput of the algorithms was measured in terms of Millions of Floating-Point Operations per Second (**MFLOPS**). This metric was calculated using the formula:

$$\text{MFLOPS} = (2 \times \text{MatrixSize}^3) / (\text{Time} \times 10^6)$$

where *MatrixSize* represents the dimension of the matrix and *Time* is the execution time in seconds.

All measurements were performed on the same machine at the computer rooms at FEUP, with an Intel Core i7 processor and the Ubuntu 22.04 operating system. The optimization flag -O2 was used on C++ algorithms. It should also be noted that each measurement was performed 3 times so as to avoid deviation of results.

## Results and Analysis

### *Simple Multiplication*

The execution time increases exponentially as the matrix size grows. This is expected due to the cubic complexity of the algorithm (Annex A2.2).

When comparing the execution time between Java and C++, we can conclude that C++ outperforms Java at a consistent rate, this is likely due to the use of the optimization flag -O2 and the lower level of memory management of C++, as it is faster than Java's garbage collector.

In this implementation, the innermost loop iterates over the columns of the second matrix (*phb*), accessing memory in a non-sequential manner. This results in poor cache utilization, as each access to *phb[k\*m\_br + j]* may cause a cache miss, especially for large matrices. This explains the high number of L1 and L2 data cache misses (DCM) observed.

### *Line Multiplication*

Lower execution times are observed for both Java and C++ when using the Line Multiplication algorithm compared to Simple Multiplication (Annex A3.2). For instance, for a 3000x3000 matrix, Line Multiplication took 16.48 seconds in C++, as opposed to 118.07 seconds for Simple Multiplication. Similarly, in Java, the execution time was reduced from 238.63 seconds to just 17.78 seconds. Notably, in Line Multiplication, the execution times for Java and C++ are very similar (Annex A3.4), unlike in Simple Multiplication, where C++ significantly outperformed Java. This improvement is likely due to Line Multiplication's optimized memory access patterns, which reduce the overhead of the Java Virtual Machine (JVM) and make Java's performance more comparable to C++.

In Line Multiplication, the order of the loops is rearranged so that the innermost loop iterates over the columns of the result matrix (*phc*) and the second matrix (*phb*). This ensures sequential memory access to *phb[k\*m\_br + j]*, improving cache utilization. As a result, the number of cache misses is significantly reduced (Annex A3.3). For example, Line Multiplication results in 6.78 billion L1 DCM and 6.30 billion L2 DCM for a 3000x3000 matrix, compared to Simple Multiplication's 50.3 billion L1 DCM and 95.8 billion L2 DCM for the same matrix size.

Since cache misses force the CPU to fetch data from the slower main memory, Line Multiplication's cache-friendly design demonstrates that reducing cache misses directly decreases execution time.

## ***Block Multiplication***

This algorithm takes the performance of both Simple Multiplication and Line Multiplication up a notch (Annex A4.2 and A4.4). For instance, for a 4096x4096 matrix with a block size of 256, Block Multiplication took 27.26 seconds, compared to 41.23 seconds for Line Multiplication and an estimated even longer time for Simple Multiplication, since this last algorithm takes about 118.07 seconds to compute the product of 3000x3000 sized matrices. Similarly, for a 10240x10240 matrix with a block size of 512, Block Multiplication completed in 428.92 seconds, while Line Multiplication took 655.97 seconds).

By dividing the matrices into blocks, the algorithm ensures that the working data *fits in* the cache, lowering the number of L1 and L2 data cache misses (DCM) (Annex A4.3). For instance, for a 4096x4096 matrix with a block size of 256, Block Multiplication resulted in 9.06 billion L1 DCM and 23.02 billion L2 DCM, compared to Line Multiplication's 17.55 billion L1 DCM and 16.25 billion L2 DCM for the same matrix size. This reduction in cache misses directly translates to faster execution times.

The choice of block size also plays a crucial role in performance. Smaller block sizes often result in higher cache misses and longer execution times, while larger block sizes strike a balance between cache utilization and computational efficiency. For example, for a 10240x10240 matrix, a block size of 512 resulted in 136.92 billion L1 DCM and 311.98 billion L2 DCM, compared to 150.47 billion L1 DCM and 512.66 billion L2 DCM for a block size of 128.

In conclusion, Block Matrix Multiplication outperforms both Simple and Line Multiplication by optimizing cache utilization through its block-oriented design. This approach reduces cache misses and execution times, demonstrating the importance of memory access patterns in high-performance computing. The results also highlight the impact of block size on performance, with larger blocks generally providing better efficiency for large matrices, as less blocks are required to cover the entire matrix and more data is loaded to the cache at once.

## ***Line Multiplication with parallel threads - First Solution***

According to the results, parallelizing the *i* loop iterations (Annex A1.4) leads to a much less acute exponential growth in both cache misses and execution time (Annexes A5.2 and A5.3).

This can be trivially explained by the fact that the work is being divided into more threads and, therefore, being executed at the same time. A speedup, in relation to the unparallelized line multiplication, of more than 6 times can be observed as up to 8 threads are being used.

However, this decrease in calculation time falls shorter as the matrix size increases, which is due to several factors, such as the increase in the overhead, the load distribution and cache limitations.

The overhead, which handles the division of the workload and the joining of the calculated values, cannot itself be parallelized. As the operations become increasingly larger, dividing

the work becomes a more resource intensive task, leading to less time spent on the matrix calculations.

Another factor is the increase in variability of the duration of execution for each thread - even though the load distribution is usually done fairly, as the execution time in each thread increases, so does the time that other threads may need to wait for.

Finally, all of the threads are sharing the same caches and therefore, when more cache accesses are needed, more collisions may occur which makes the threads have to wait longer for those accesses.

With all of this in mind, we can conclude that, even though the parallelization of the calculation clearly leads to an improved performance, the results demonstrate some details in the way it functions that result in some losses in speedup in exponential complexity work.

### *Line Multiplication with parallel threads - Second Solution*

In this solution, the  $j$  iterations are executed in parallel, i.e., less work is being parallelized more often. As the results show, this is clearly the less optimal scenario of the two when it comes to using multiple threads. The speedup in relation to non parallel line multiplication (Annex A6.1), even though positive in most cases, far underperforms when compared to the first solution. The number of threads may be the same, but the MFLOPS are noticeably less, which leads to a lower efficiency (Annexes A6.6 and A6.7).

This is a poor setup for parallelizing work as the factors that take part in reducing the speedup in the first solution are even more evident. The overhead work is much more complex as the execution is being distributed much more frequently, which increases the amount of sequential work that needs to be done. This can be observed as, in low matrix sizes, the speedup is negative (Annex A6.1) in comparison to the single thread algorithm. If we take into consideration the use of several threads, this process seems even less helpful as a lot of cores are being kept busy for minimal gains.

## **Conclusions**

This project helped us achieve a better understanding of how several cores may be used to improve a process's performance and how to better decide what sections of the code should be parallelized. The overhead was shown to affect execution time significantly and reducing it is important to develop a good multi-thread program. It has also shown us that the way we design the matrix multiplication algorithm, even when using a single thread, may affect the performance, as evidenced by the greater performance of the block multiplication.

# Annexes

## A1 - Code Snippets

### A1.1 - Simple Multiplication

```
for(i=0; i<m_ar; i++)
{
    for( j=0; j<m_br; j++)
    {
        temp = 0;
        for( k=0; k<m_ar; k++)
        {
            temp += pha[i*m_ar+k] * phb[k*m_br+j];
        }
        phc[i*m_ar+j]=temp;
    }
}
```

### A1.2 - Line Multiplication

```
for(i=0; i<m_ar; i++){
    for( k=0; k<m_ar; k++ ){
        for( j=0; j<m_br; j++){
            phc[i*m_ar+j] += pha[i*m_ar+k] * phb[k*m_br+j];
        }
    }
}
```

### A1.3 - Block Multiplication

```
for (ii = 0; ii < m_ar; ii += bkSize) {
    for (kk = 0; kk < m_br; kk += bkSize) {
        for (jj = 0; jj < m_ar; jj += bkSize) {
            for (i = ii; i < min(ii + bkSize, m_ar); i++) {
                for (k = kk; k < min(kk + bkSize, m_br); k++) {
                    for (j = jj; j < min(jj + bkSize, m_ar); j++) {
                        phc[i * m_ar + j] += pha[i * m_ar + k] *
phb[k * m_br + j];
                    }
                }
            }
        }
    }
}
```



#### A1.4 - Line Multiplication with parallel threads - First Solution

```
#pragma omp parallel for
    for(i=0; i<m_ar; i++){
        for( k=0; k<m_ar; k++){
            for( j=0; j<m_br; j++){
                phc[i*m_ar+j] += pha[i*m_ar+k] * phb[k*m_br+j];
            }
        }
    }
```

#### A1.5 - Line Multiplication with parallel threads - Second Solution

```
#pragma omp parallel
    for(i=0; i<m_ar; i++){
        for( k=0; k<m_ar; k++){
            #pragma omp parallel for
            for( j=0; j<m_br; j++){
                {
                    phc[i*m_ar+j] += pha[i*m_ar+k] * phb[k*m_br+j];
                }
            }
        }
    }
```

## A2 - Simple Multiplication - Measurements and Graphs

### A2.1 - Measurements Table

Simple Multiplication Measurements - C++ vs Java

| Size | Average Time (s) - C++ | Average Time (s) - Java | Average L1 DCM | Average L2 DCM |
|------|------------------------|-------------------------|----------------|----------------|
| 600  | 0.19                   | 0.27                    | 244,695,917    | 39,368,746     |
| 1000 | 1.35                   | 13.34                   | 1,224,484,542  | 295,071,267    |
| 1400 | 3.64                   | 14.96                   | 3,465,951,803  | 1,292,634,886  |
| 1800 | 17.66                  | 36.07                   | 9,085,822,351  | 4,727,840,269  |
| 2200 | 38.37                  | 77.08                   | 17,627,555,173 | 18,284,610,671 |
| 2600 | 70.76                  | 121.08                  | 30,904,351,409 | 50,155,807,448 |
| 3000 | 118.07                 | 238.63                  | 50,289,207,683 | 95,779,490,469 |

### A2.2 - C++ vs Java Execution Time vs Matrix Size




## A3 - Line Multiplication - Measurements and Graphs

### A3.1 - Measurements Table


Line Multiplication Measurements - C++ vs Java

| Size  | Average Time (s) - C++ | Average L1 DCM - C++ | Average L2 DCM - C++ | Average Time (s) - Java |
|-------|------------------------|----------------------|----------------------|-------------------------|
| 600   | 0.115                  | 27,111,997           | 57,787,192           | 0.12                    |
| 1000  | 0.475                  | 125,734,561          | 267,747,076          | 0.55                    |
| 1400  | 1.583                  | 346,116,682          | 708,816,444          | 1.65                    |
| 1800  | 3.557                  | 745,271,243          | 1,439,900,904        | 3.57                    |
| 2200  | 6.473                  | 2,073,886,158        | 2,561,167,064        | 6.57                    |
| 2600  | 10.819                 | 4,412,845,122        | 4,198,690,629        | 11.07                   |
| 3000  | 16.478                 | 6,780,698,787        | 6,296,071,144        | 17.78                   |
| 4096  | 41.229                 | 17,553,265,709       | 16,249,690,528       | nan                     |
| 6144  | 139.299                | 59,151,207,256       | 54,420,319,247       | nan                     |
| 8192  | 336.908                | 140,290,893,728      | 130,816,392,740      | nan                     |
| 10240 | 655.973                | 273,727,875,965      | 256,199,650,326      | nan                     |


### A3.2 - C++ Simple vs. Line Multiplication: Time Comparison



### A3.3 - C++ Simple vs. Line Multiplication: Cache Misses Comparison



### A3.4 - C++ vs. Java Line Multiplication: Time Comparison



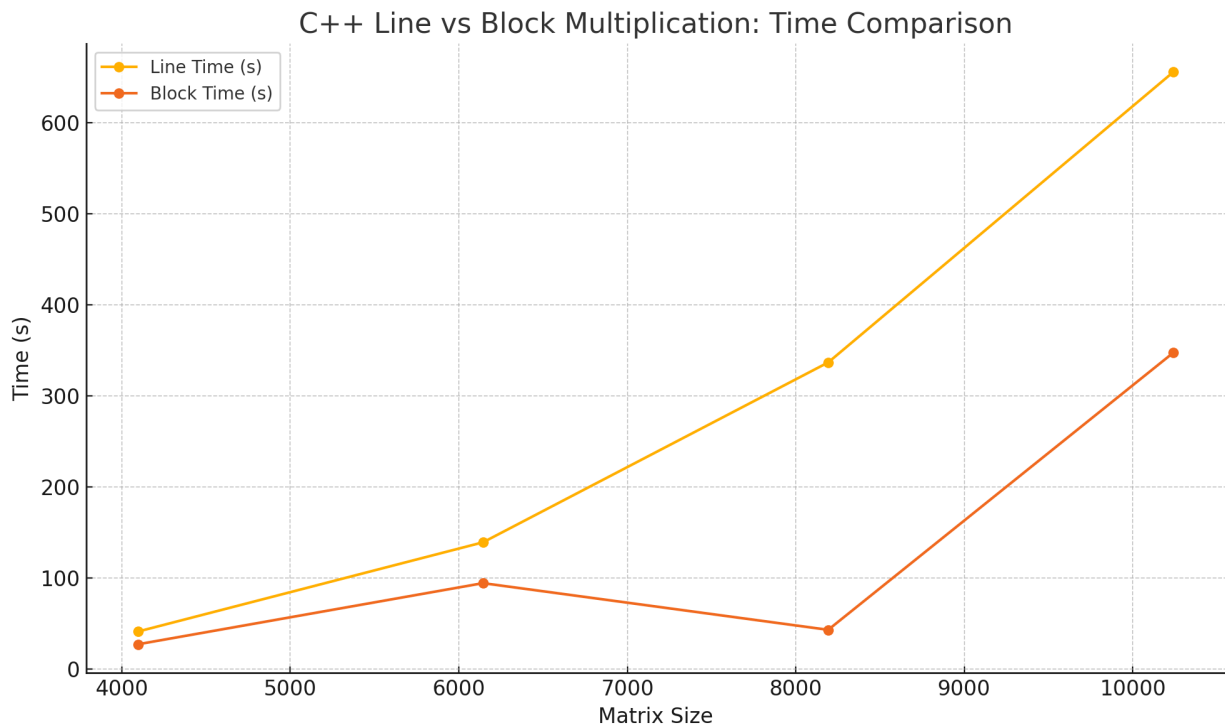
## A4 - Block Multiplication - Measurements and Graphs

### A4.1 - Measurements Table

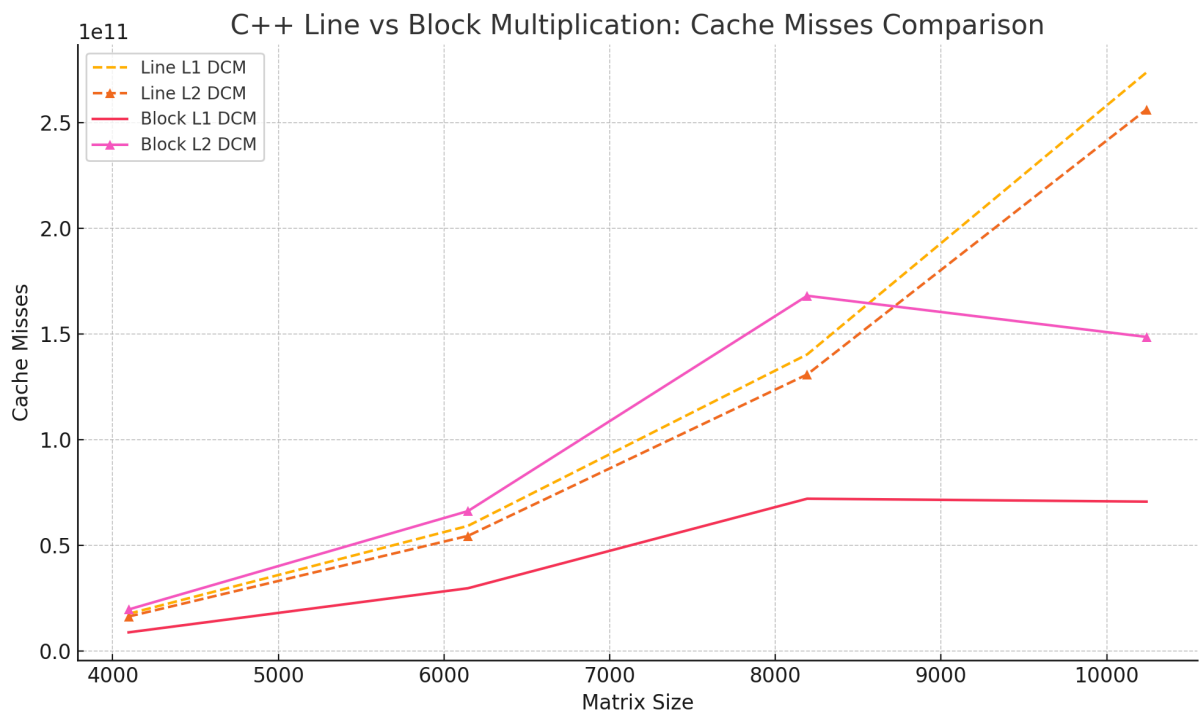
Block Multiplication Measurements - C++

| Size  | Block Size | Average Time (s) | Average L1 DCM  | Average L2 DCM  |
|-------|------------|------------------|-----------------|-----------------|
| 4096  | 128        | 31.34            | 9,659,705,054   | 32,416,819,339  |
| 4096  | 256        | 27.264           | 9,062,584,580   | 23,017,786,804  |
| 4096  | 512        | 28.516           | 8,768,735,120   | 19,594,615,939  |
| 6144  | 128        | 108.801          | 32,486,518,058  | 109,657,143,152 |
| 6144  | 256        | 94.431           | 30,547,622,618  | 77,903,305,841  |
| 6144  | 512        | 96.499           | 29,641,740,553  | 66,145,750,807  |
| 8192  | 128        | 303.677          | 74,278,262,468  | 256,665,253,340 |
| 8192  | 256        | 43.166           | 72,020,268,392  | 168,037,026,742 |
| 8192  | 512        | 347.631          | 70,629,754,365  | 148,556,814,500 |
| 10240 | 128        | 505.63           | 150,474,644,703 | 512,657,128,442 |
| 10240 | 256        | 431.71           | 141,972,681,471 | 356,982,368,547 |
| 10240 | 512        | 428.924          | 136,921,824,089 | 311,978,693,469 |

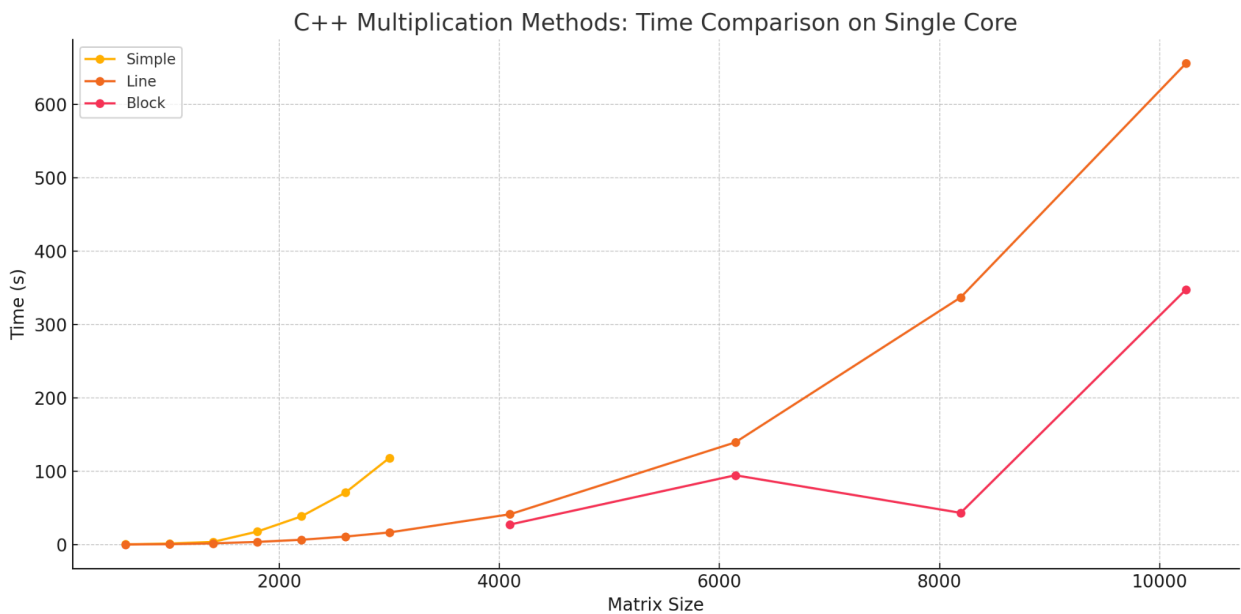
### A4.2 - C++ Line vs. Block Multiplication - Time Comparison



#### A4.3 - C++ Line vs. Block Multiplication - Cache Misses Comparison



#### A4.4 - C++ Simple vs. Line vs. Block Multiplication - Time Comparison




## A5 - Line Multiplication with Parallel Threads - First Solution - Measurements and Graphs

### A5.1 - Measurements Table

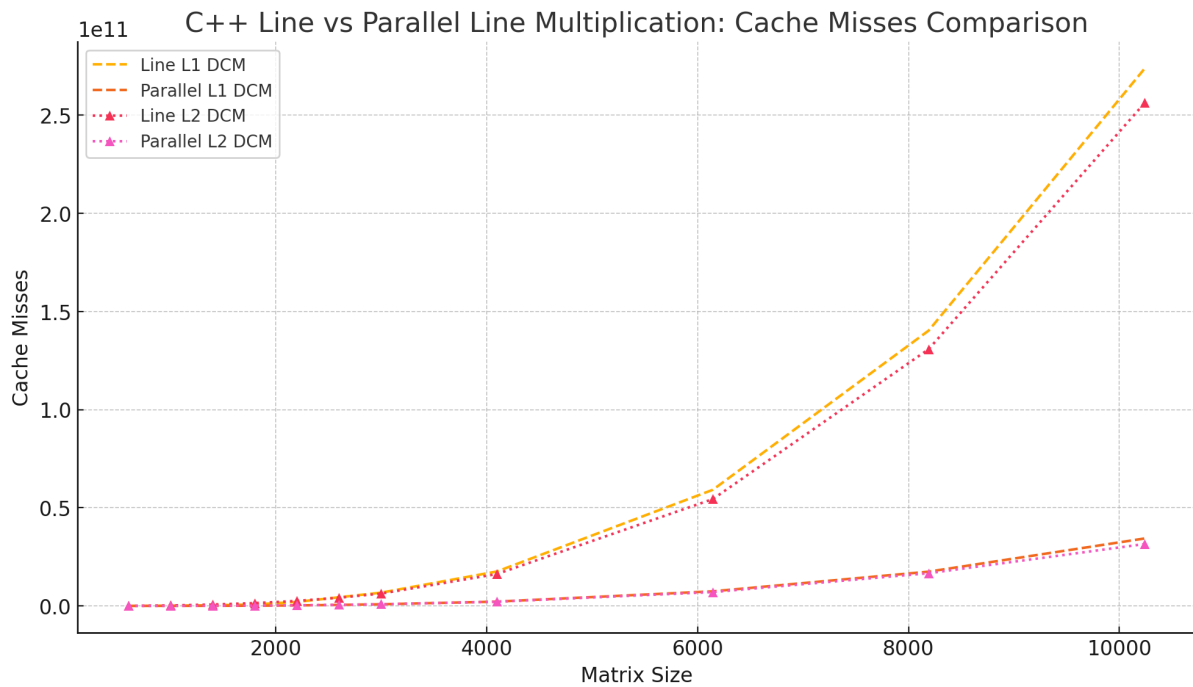
Parallel Multiplication Measurements (First Solution)

| Size  | Average Time (s) | Average L1 DCM | Average L2 DCM | Mflops   | Speedup | Efficiency |
|-------|------------------|----------------|----------------|----------|---------|------------|
| 600   | 0.018            | 3,390,655      | 7,174,854      | 24000.0  | 6.389   | 79.86%     |
| 1000  | 0.078            | 15,713,594     | 32,649,054     | 25641.03 | 6.09    | 76.12%     |
| 1400  | 0.244            | 43,437,512     | 87,413,699     | 22491.8  | 6.488   | 81.10%     |
| 1800  | 0.552            | 93,484,317     | 184,974,373    | 21130.43 | 6.444   | 80.55%     |
| 2200  | 1.065            | 258,616,185    | 331,133,590    | 19996.24 | 6.078   | 75.97%     |
| 2600  | 1.76             | 549,797,831    | 548,448,835    | 19972.73 | 6.147   | 76.84%     |
| 3000  | 2.816            | 845,412,912    | 836,080,513    | 19176.14 | 5.852   | 73.14%     |
| 4096  | 7.514            | 2,197,529,677  | 2,132,961,687  | 18291.05 | 5.487   | 68.59%     |
| 6144  | 29.642           | 7,431,951,547  | 7,049,828,317  | 15648.62 | 4.699   | 58.74%     |
| 8192  | 82.798           | 17,459,326,625 | 16,672,832,914 | 13279.45 | 4.069   | 50.86%     |
| 10240 | 178.491          | 34,362,758,953 | 31,519,192,325 | 12031.33 | 3.675   | 45.94%     |

### A5.2 - Line vs Parallel Line 1: Time Comparison



### A5.3 - Line vs Parallel Line 1: Cache Misses Comparison



## A6 - Line Multiplication with Parallel Threads - Second Solution - Measurements and Graphs

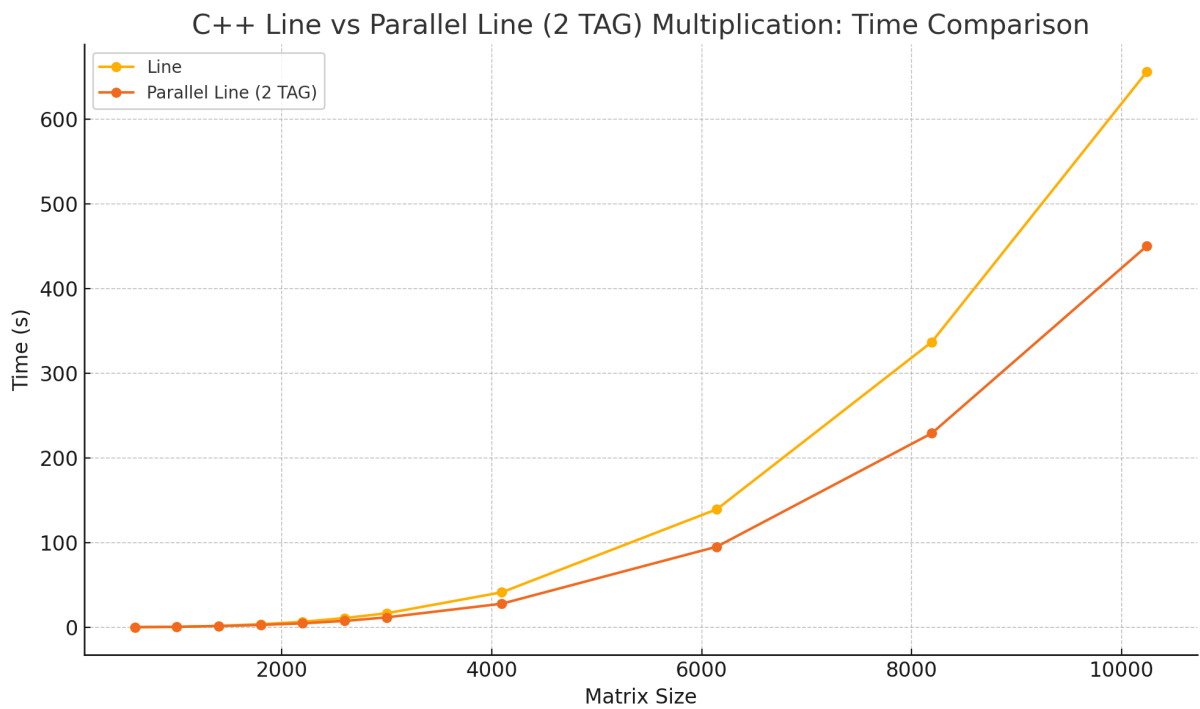
### A6.1 - Measurements Table

Parallel Multiplication Measurements (Second Solution)

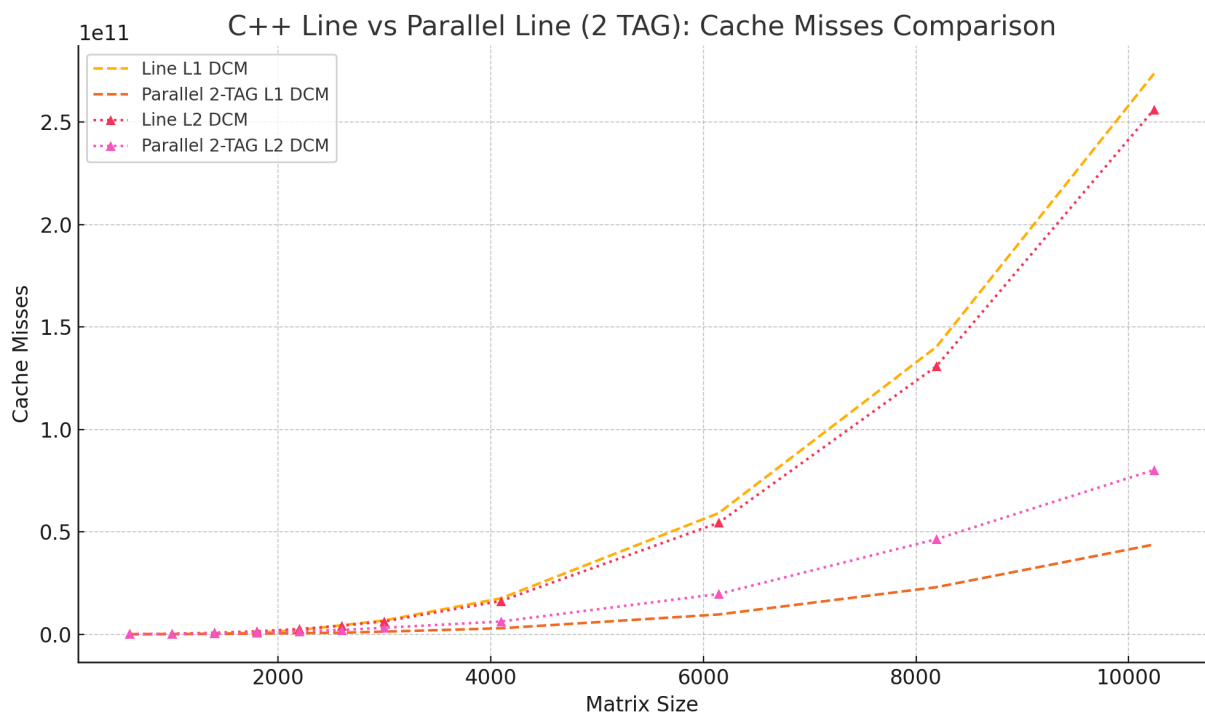
| Size  | Average Time (s) | Average L1 DCM | Average L2 DCM | Mflops  | Speedup | Efficiency |
|-------|------------------|----------------|----------------|---------|---------|------------|
| 600   | 0.121            | 12,687,581     | 43,001,109     | 3570.25 | 0.95    | 11.88%     |
| 1000  | 0.513            | 57,322,579     | 186,063,316    | 3898.64 | 0.926   | 11.57%     |
| 1400  | 1.408            | 149,451,412    | 464,926,561    | 3897.73 | 1.124   | 14.05%     |
| 1800  | 2.703            | 309,991,421    | 914,588,639    | 4315.21 | 1.316   | 16.45%     |
| 2200  | 4.629            | 514,954,420    | 1,457,204,131  | 4600.56 | 1.398   | 17.48%     |
| 2600  | 7.448            | 782,717,803    | 2,138,071,919  | 4719.66 | 1.453   | 18.16%     |
| 3000  | 11.662           | 1,265,554,735  | 3,173,222,781  | 4630.42 | 1.413   | 17.66%     |
| 4096  | 27.762           | 2,927,306,005  | 6,236,507,766  | 4950.61 | 1.485   | 18.56%     |
| 6144  | 95.12            | 9,667,315,438  | 19,660,774,794 | 4876.54 | 1.464   | 18.31%     |
| 8192  | 229.101          | 22,984,441,740 | 46,412,806,000 | 4799.24 | 1.471   | 18.38%     |
| 10240 | 450.163          | 43,823,951,312 | 80,218,070,755 | 4770.46 | 1.457   | 18.21%     |




### A6.2 - Line vs. Parallel Line 2: Time Comparison



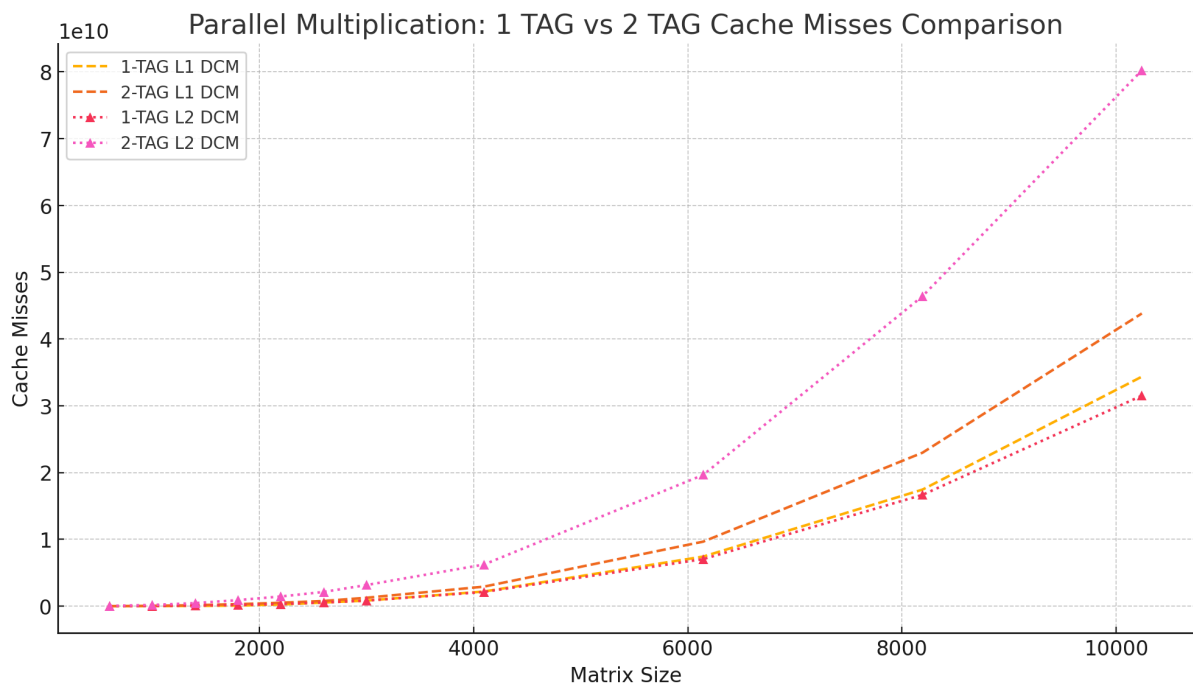
### A6.3 - Line vs. Parallel Line 2: Cache Misses Comparison



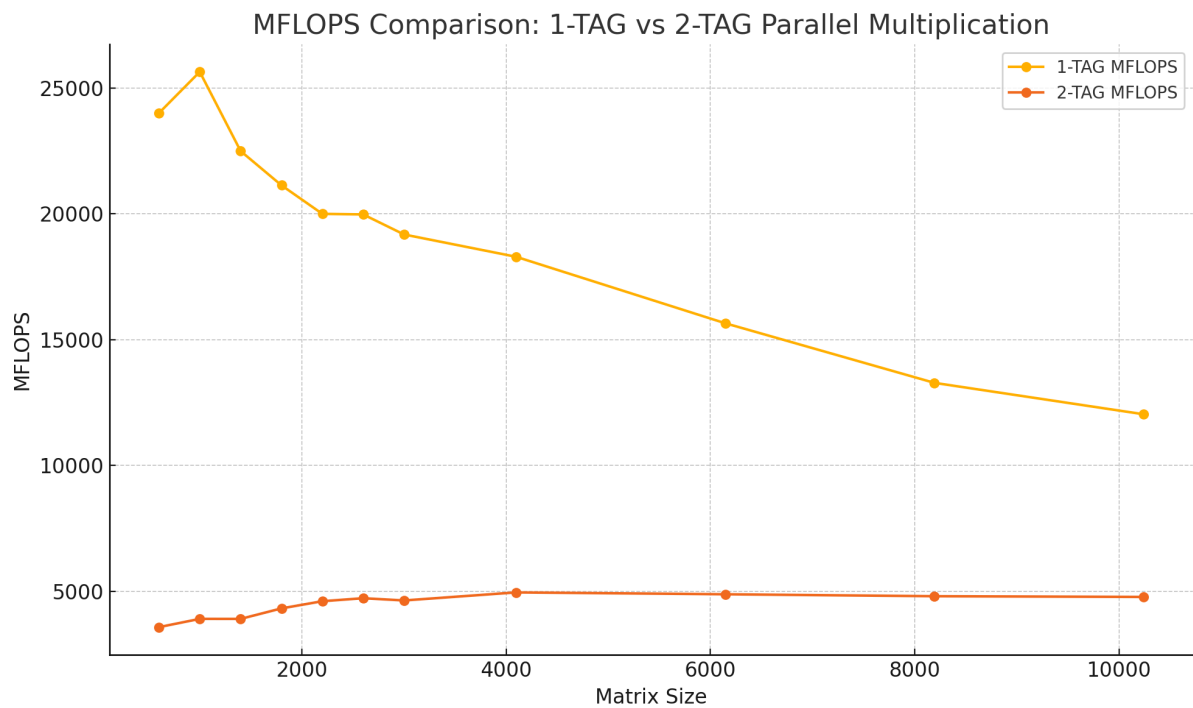
#### A6.4 - Parallel 1 vs. Parallel 2 - Time Comparison



### A6.5 - Parallel 1 vs. Parallel 2 - Cache Misses Comparison



### A6.6 - Parallel 1 vs. Parallel 2 - MFLOPS Comparison



### A6.7 - Parallel 1 vs. Parallel 2 - Efficiency Comparison

